

Week 3 Quiz
Graded Assignment

← Back

English ▾ Due Jun 8, 10:59 PM CST

Your grade: 100%

Your latest: 100% • Your highest: 100%

To pass you need at least 80%. We keep your highest score.

Next item →

1. When stacking LSTMs, how do you instruct an LSTM to feed the next one in the sequence?

1 / 1 point

- ☐ Do nothing, TensorFlow handles this automatically
- ☐ Ensure that return_sequences is set to True on all units
- ☒ Ensure that return_sequences is set to True only on units that feed to another LSTM
- ☐ Ensure that they have the same number of units

✔ Correct
Correct!

2. How does an LSTM help understand meaning when words that qualify each other aren't necessarily beside each other in a sentence?

1 / 1 point

- ☐ They load all words into a cell state
- ☐ They shuffle the words randomly
- ☐ They don't
- ☒ Values from earlier words can be carried to later ones via a cell state

✔ Correct
Correct!

3. What's the best way to avoid overfitting in NLP datasets?

1 / 1 point

- ☐ Use LSTMs
- ☐ Use GRUs
- ☐ Use Conv1D
- ☒ None of the above

✔ Correct
Correct!

4. What keras layer type allows LSTMs to look forward and backward in a sentence?

1 / 1 point

- ☐ Bothdirection
- ☐ Unilateral
- ☒ Bidirectional
- ☐ Bilateral

✔ Correct
Correct!

5. Why does sequence make a large difference when determining semantics of language?

1 / 1 point

- ☐ It doesn't
- ☐ Because the order in which words appear dictate their meaning
- ☒ Because the order in which words appear dictate their impact on the meaning of the sentence
- ☐ Because the order of words doesn't matter

✔ Correct
Correct!

6. How do Recurrent Neural Networks help you understand the impact of sequence on meaning?

1 / 1 point

- ☒ They carry meaning from one cell to the next
- ☐ They shuffle the words evenly
- ☐ They don't
- ☐ They look at the whole sentence at a time

✔ Correct
That's right!

7. What's the output shape of a bidirectional LSTM layer with 64 units?

1 / 1 point

- ☐ (128,1)
- ☐ (128,None)
- ☒ (None, 128)
- ☐ (None, 64)

✔ Correct
That's right!

8. If a sentence has 120 tokens in it, and a Conv1D with 128 filters with a Kernel size of 5 is passed over it, what's the output shape?

1 / 1 point

- ☒ (None, 116, 128)
- ☐ (None, 120, 128)
- ☐ (None, 120, 124)
- ☐ (None, 116, 124)